

Bachelor Thesis

**Automata Learning at Scale: Evaluation of
Query Parallelization Strategies in the
Context of Clustered Systems**

Julien Saevecke

November 10, 2022

Reviewer:

Prof. Dr. Bernhard Steffen

M.Sc. Julius Alexander Bainsczyk

Technische Universität Dortmund

Fakultät für Informatik

Programmiersysteme (LS-5)

<http://ls5-www.cs.tu-dortmund.de>

Contents

1	Introduction	1
1.1	Related Work	2
1.2	Goals	3
1.3	Structure	4
2	Preliminaries	5
2.1	Active Automata Learning	5
2.1.1	LearnLib	6
2.2	Mealy Machine	7
2.3	Virtualization Technologies	9
2.4	Container Orchestration	10
3	Automata Learning of Clustered Systems	13
3.1	Membership Query Distribution	13
3.2	Challenges Of Real-Life Applications	16
3.3	Orchestrating Containerized SULs	16
3.4	Strategies for Scaling SULs	18
3.4.1	Horizontal Pod Autoscaler	19
3.4.2	Kubernetes Event-driven Autoscaling	20
4	Evaluation Method	23
4.1	Experiment: Simple Coffee Machine	23
4.2	SUL Containerization	24
4.3	Simulating Real-Life Systems	26
4.4	Data Acquisition	26
4.5	Learning Setups	27
5	Results	29
5.1	Architecture Overhead	29
5.2	Strategy: Constant Number Of SULs	31
5.2.1	Introducing Processing Time Delay	32

5.3 Strategy: Auto-Scaling	33
5.4 Evaluation	34
6 Conclusion	37
6.1 Future Research	38
A Index Of Abbreviations	39
Index Of Abbreviations	41
Index Of Figures	41
Index Of Algorithms	43
Bibliography	50
Declaration	50

Chapter 1

Introduction

These days, computer software is used in many objects in the environment. It begins with the apparent computer helping us in our daily life, to crucial medical equipment for preserving or saving lives, and ends with self-driving cars for transporting goods or people. Regrettably, even the simplest software is prone to bugs expressing themselves differently. A bug in software can express itself as a simple button on a website not working. However, it can also lead to catastrophic financial damage or, even worse, to loss of life.[68] Thus, it is important to prevent software failure or bugs that can influence lives in a bad way or even endanger them. Consequently, between 30 and 60 percent of the software development process consists of testing.[66]

One effective approach to prevent or reduce software failures and raise software quality is Model-Based Testing (**MBT**). In this approach, a model, for example, a Finite-State Machine (**FSM**), is used to represent the desired behavior of a System Under Test (**SUT**). With the help of the model, test cases can be generated to test whether the implementation of the **SUT** is working correctly. In **MBT**, creating the model can be time-consuming and cost-inefficient when done manually and vulnerable to errors, especially regarding complex real-life systems.

Active automata learning is a technique to infer a behavioral model representing the actual behavior by executing test queries on the **SUT** (in this context, further called System Under Learning (**SUL**)). A framework popular for setting up active automata learning is LearnLib[18]. However, this technique has two major concerns, which are further emphasized in the context of real-life systems.

The first primary concern is that active automata learning is a rather costly technique because gaining sufficient knowledge for inferring a behavioral model of the target system is glued to a significant number of test queries sent to the **SUL**. Depending on the system's complexity, responding to such a test query may induce high latency operations and take a long time to process. Consequently, learning behavioral models of such a system may take hours, days, or even months.[56]

The second concern is that active automata learning requires test query independence to learn correctly. Meaning no test query affects another test query. Therefore, the requirement on the target system is to offer reset functionality or any other abstractions guaranteeing that no query impacts the other. However, in the context of real-life systems, this requirement is often not fulfilled because the system may consist of, e.g., of multiple services and persistent storage to maintain any state.[46]

These two concerns were why inferring behavioral models of real-life systems was not practicable. However, in 2011 Next Generation LearnLib (**NGLL**)(further referred to as LearnLib for simplicity), with the focus on offering various features making the learning of real-life systems more feasible, such as reset mechanisms and abstraction/refinement techniques, was released. Furthermore, it is implemented in an extensible modular fashion, allowing active automata learning on different system architectures, e.g., a distributed system, thus opening up the possibility of more flexible learning setups. [47]

This change led to more research focusing on real-life applications, such as making active automata learning less costly and introducing so-called mappers[59] as a communication interface to real-life systems. Consequently, active automata learning was successfully used to learn various behavioral models. In addition, the tool Automata Learning Experience (**ALEX**)[2] for inferring models of web applications and JSON-based web services was released and has been actively developed since 2015.

Despite progress in reducing the total learning time and enabling the learning of real-life systems, active automata learning is still a heavily weighted process. Therefore, this work proposes and evaluates an architecture that enables scalable active automata learning with accommodations to real-life systems.

1.1 Related Work

Performing active automata learning is a heavily weighted process, especially in the context of real-life systems. A learning process may take a long time to learn a behavioral model of the target system. Therefore, often not feasible to utilize in **MBT** in, e.g., agile or time-constrained environments. Thus, existing studies examined various methods to reduce the total execution time of the active automata learning process, such as: optimizing equivalence queries, using learn algorithms[57, 56], reusing states(removing duplicated work)[30, 37], and abstraction layers (e.g., using equivalence classes)[25].

In [37], active automata learning is evaluated in a context combining multi-threading and reusing states and is compared to the usually single-threaded approach taken in the research mentioned above. This approach runs multiple **SULs** simultaneously as threads on a single machine. It divides the Membership Query (**MQ**)s, which are sent in the learning stage, and Equivalence Query (**EQ**)s before being sent into numerous batches and distributed to the running **SULs** to balance the load.

Additionally, [46] examined active automata learning with a primary focus on real-life applications, regarding the challenges of learning them and which steps can be taken to make learning such systems more practicable. This work follows the same approach as [37] by parallelizing and distributing queries to multiple SULs on a distributed system and compares different learning algorithms regarding the query batch size affecting the total execution time in this context.

However, no in-depth evaluation was made using active automata learning on real-life applications performed on a distributed system. Thus, this work tries to close this gap by tackling the challenges of active automata learning regarding real-life applications on distributed systems and evaluating such a setup's performance and practicability.

1.2 Goals

The approach in this research complements the scale-able aspect of the method taken in [37] and shifts the context, similar to the cloud computing approach [46], from the active automata learning process running on a single physical machine into a clustered system. A clustered system contains a group of at least two physical and virtual machines, interconnected with each other in a way that they can act like a single system. Operating active automata learning on a clustered system through parallelizing and distributing the processing of queries to several SUL instances can improve the execution time to learn a behavioral model notably and is not limited by the resources of a single machine.

However, a clustered system introduces management efforts to the general setup to accomplish that. Still, utilizing virtualization technology and orchestration tools, it can provide and manage resources on demand, therefore, adapt to ever-changing workloads by scaling resources up and down while running and being resilient. Therefore this work has the following primary goal:

Construct an architecture utilizing a clustered system to perform scale-able active automata learning

Consequently, this research proposes an architecture to perform scale-able active automata learning. Utilizing a clustered system and technologies enables the simultaneous processing of test queries distributed to multiple SULs at once. Furthermore, run-time auto-scaling strategies are used to efficiently utilize the available cluster resources and compared to an approach without scaling. The target SULs are adjusted to accommodate real-life systems using processing delays and restarting them to ensure test query independence.

Resulting in an architecture that can learn all kinds of SULs without offering additional functionality such as reset mechanisms. Furthermore, this approach is in-depth evaluated

regarding practicability, scale-ability, and limitations, focusing on real-life systems. Thus, closing the research gap mentioned above.

1.3 Structure

On that account, this work is structured in the following way. In the first part ([chapter 2](#)), the fundamentals of active automata learning and the process of learning a behavioral model of a [SUL](#) are covered. Furthermore, it is essential to cover mealy machines and the tool LearnLib in this context. Finally, the technologies used to set up active automata learning in a clustered system are explained to complete the fundamentals.

In the second part ([chapter 3](#) and [chapter 4](#)), the proposed architecture of the clustered automata learn setup is explained, accompanied by an explanation of how experiments are set up, run, and measured for later evaluation.

Before last, the results of the run learning setups are presented, analyzed, and evaluated ([chapter 5](#)). Insights about practicability, limitations, and performance concerning the proposed architecture's scale-ability are disclosed in this part.

The last part ([chapter 6](#)) concludes this research by determining whether the goals stated in [section 1.2](#) were accomplished. In addition, ideas for further research focusing on possible improvements or extensions to the proposed architecture may lead to performance improvements or enable the architecture to learn real-life systems.

Chapter 2

Preliminaries

Before explaining how proposed architecture can be implemented to utilize the underlying clustered system to perform scaleable active automata learning, it is essential to introduce the fundamentals of the technologies used to make that happen. In the same breath, the tool LearnLib is examined regarding how its structure and features can help to move the active automata learning process to a clustered system.

2.1 Active Automata Learning

Active Automata Learning is a process that automatically infers formal behavioral models of, e.g., black box systems through testing. [61] Furthermore, inferred behavioral models enrich the development of systems by using, e.g., quality assurance.

The indicated process adheres to the Minimally Adequate Teacher (MAT) model and is comparable to the analogy of a relationship between a learner and a teacher (further called SUL). In the context of active automata learning, the learner has to figure out an unknown formal language L known by the SUL by requesting answers from the SUL on specific questions (further called queries)[50, 46]:

1. A Membership query (MQ) is a sequence $w \in \Sigma^*$, whereas Σ is given out by the teacher. Based on w , that gets processed by the SUL, leading to an answer. The learner uses this answer to construct a hypothesis model of the unknown target L
2. A Equivalence query (EQ) is defined like a MQ. However, the difference is that this query is used to compare the current constructed hypothesis model of the learner with the actual behavioral model of L known by the SUL. This comparison may result in inequality, and a counterexample is provided, proofing the dissimilarity

Using these queries, the learner follows an alternating two-stage strategy to determine the L . At first, the MQs are used to construct the hypothesis model of the unknown L , followed by testing the current hypothesis model through EQs may result in a counterexample. A

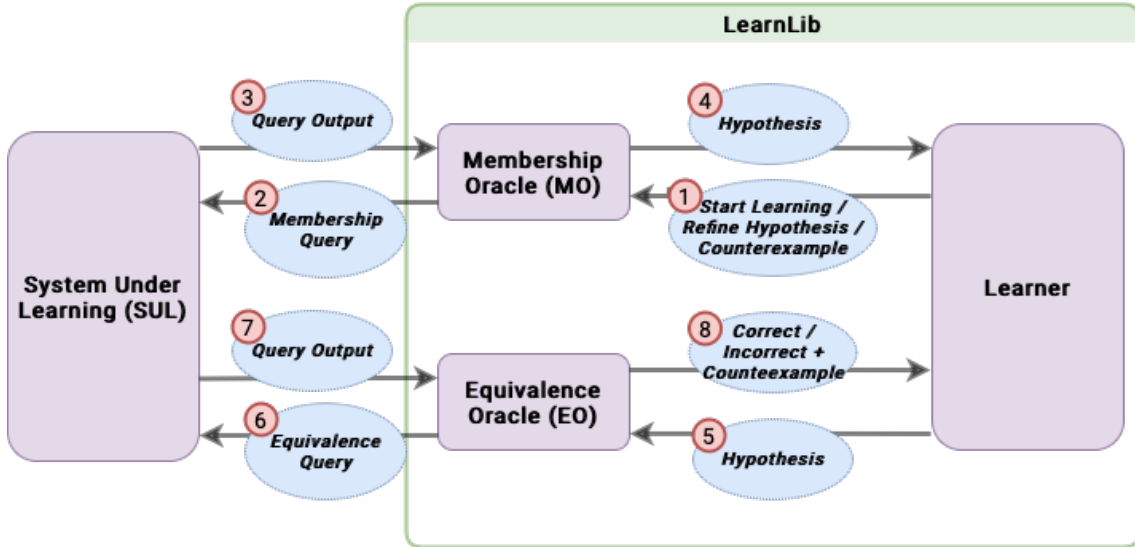


Figure 2.1: Learn Process With LearnLib Based On [37]

refined hypothesis model can be constructed using the counterexample, which may result in new **MQs** and **EQs**. As stated, these two stages take turns until **EQs** no longer produce counterexamples. At this point, the hypothesis model is equal to the target L .

However, this work focuses on real-life systems, usually black-box systems that introduce some challenges. Firstly, the hypothesis model can not be directly compared to the target model of the **SUL**; thus, the **EQs** have to be approximated. Consequently, this may result in a higher **MQ** and **EQ** count than learning a non-black-box system.[62]

Secondly, **MQs** need to be independent of each other - meaning the **SUL** in question has to be resettable to its initial state before each **MQ**. This requirement can be problematic because real-life systems often rely on an internal state stored in, e.g., databases and may contain multiple services. Such systems may not offer functionality to reset the whole system to its initial state without restarting it. [40]

At last, simple input alphabets often do not suffice to communicate with real-life systems; thus, a so-called "mapper" is used, representing an intermediate layer between the learner and the **SUL**, which translates **MQs**, e.g., to method invocations.[59, 39] The first two challenges are tackled in this work.

There is a choice to be made when it comes to libraries enabling automata learning adhering to the explained **MAT** model: LearnLib[18] or libalf[20]. This work uses LearnLib as the library of choice. The reason is the faster performance compared to the other available choices and its flexible modular structure.[41, 19]

2.1.1 LearnLib

LearnLib is an open-source framework built in Java[13] specializing in active automata learning. It provides a modular design enabling flexible and extensible learning configura-

tions and allows for capitalization of specific properties of the target **SUL**. [50] Additionally, the design structure accommodates using different infrastructures, e.g., distributed systems while maintaining a good performance [41, 46, 37].

Further, it provides various learning algorithms, equivalence approximation strategies, and filters such as caches to reduce the total number of queries. [46] Thus, LearnLib was used to learn models of various implementations. Some of these implementations are: TCP[42], SSH[63], bank cards[24] and printer software[56].

A more in-depth active automata learning process, implemented with LearnLib and adhering to the **MAT** model, is illustrated in **Figure 2.1**. It can be separated into two parts. The first part is the target system that should be learned, namely the **SUL**. The second is the learning process implemented in LearnLib, which consists of the following components:

- Membership Oracle (**MO**) is a simple interface to a **SUL**, acting as an in-between layer - handling the communication between the **SUL** and the learner. It sends the **MQs** generated by the learning algorithm of the learner to the **SUL**. It sends back the responses to the learner, therefore helping in creating/refining the hypothesis model.
- Equivalence Oracle (**EO**) produces the already mentioned **EQs**. The structure is the same as the **MQs**. The difference is that these queries are sent to the target **SUL** by the **EO** to check the validity of the constructed hypothesis model. A chosen algorithm provided by LearnLib might approximate these queries in the case of black-box systems

To conclude, moving the active automata learning process into the context of a clustered system is possible due to the extensible framework of LearnLib. Thus, LearnLib and its extensible framework are utilized in this work to implement, e.g., a custom **MO** to adhere to the unique requirements of a clustered system (explained in **section 3.1**).

2.2 Mealy Machine

LearnLib offers various ways to implement and learn finite deterministic automata, e.g., Mealy machines. [41] In this work, it is crucial to understand how Mealy machines work because the behavior of the **SUL** used to evaluate the architecture is represented as Mealy machines.

Mealy machines are a variant of a **FSM**, where an output alphabet accompanies the input alphabet. The output is determined by the current input and the current state of the Mealy machine. Mealy machines differentiate themselves from other **FSMs** by having a transition function defined for all input symbols resulting in the output to a sequence of inputs being deterministic. A state diagram can represent Mealy machines.




Figure 2.2: A simple coffee machine as a Mealy machine [60]

2.2.1 Definition. The formal definition of a Mealy machine is a six-tuple $M = (S, s_0, \Sigma, \Omega, \delta, \lambda)$ where

- S is a finite nonempty set of states
- $s_0 \in S$ is the initial state,
- Σ is a finite input alphabet,
- Ω is a finite output alphabet,
- $\delta : S \times \Sigma \rightarrow S$ is the transition function, and
- $\lambda : S \times \Sigma \rightarrow \Omega$ is the output function.

The example Mealy machine state diagram shown in [Figure 2.2](#) represents a simple coffee machine. This state diagram defines a coffee machine with four different interactions: fill in water (water), put in a coffee pod (pod), start coffee making process (button), and clean the machine (clean). These interactions can occur in arbitrary order and based on the sequence of interactions.

The following responses are possible: valid interaction (✓), invalid interaction (★), and making coffee (☕). Nevertheless, it is required to have enough water and a coffee pod in the coffee machine before pressing the button to make a cup of coffee successfully (e.g., $\{clean, pod, water, coffee\} = ☕$). All other sequences in-between either lead to an invalid action (★) or are missing further interactions (✓) to make a cup of coffee.

The behavior of the simple coffee machine explained above is learned with various configurations to evaluate the proposed architecture of this work. More details on the specifications of the coffee machine are covered in [section 4.1](#).

2.3 Virtualization Technologies

It is essential to utilize the cluster resources efficiently to effectively perform scalable automata learning to enable as many [SULs](#) as possible for parallelization on the clustered system. This can be achieved by having [SULs](#) that are easy to distribute, replicate, and fast to deploy and run. Therefore, it is essential that [SULs](#) utilize as little as possible Central Processing Unit ([CPU](#)) and memory to run a high quantity of [SULs](#) simultaneously to process a high count of queries in parallel. Furthermore, it is required that deploying a new [SUL](#) or terminating an existing [SUL](#) replica is possible while a learning process is still running to enable scaling on-demand at run-time. For this purpose, the use of virtualization technology is vital.

Virtualization generally is a technology that partitions physical system resources to create an abstraction of one or many isolated environments. [\[52, 34\]](#) Virtualization is used to make resource utilization more efficient and provide more flexibility. [\[31\]](#) This technology can be differentiated into multiple types of virtualization. The most used techniques are hypervisor-based and Operating System ([OS](#))-level virtualization. [\[27, 28\]](#)

Virtualization on the hardware level is accomplished through a hypervisor implemented directly on top of the host machine's hardware. The hypervisor is the intermediate software layer that can launch multiple completely isolated guest [OS](#). One such system is referred to as a Virtual Machine ([VM](#)) or as full virtualization. Each [VM](#) gets its resources through the hypervisor, acting as they run directly on the hardware. [\[52, 34\]](#) This technique brings an overhead because each [VM](#) runs its [OS](#) and file system. This overhead is why a [VM](#) is slow to launch[\[28, 51\]](#) and should be used in scenarios where complete isolation and security, ease of management, and customization are the primary features in need.[\[53, 34\]](#)

The [OS](#)-level virtualization technique is also known as Container-based operating system ([COS](#)) virtualization[\[58\]](#) or containerization[\[65\]](#). Instead of virtualizing the underlying hardware of the host machine with an intermediate software layer like a hypervisor, the container-based virtualization approach virtualizes the [OS](#) kernel without a need for an additional software layer by creating multiple isolated user-space environments by dividing the physical resources of the host machine usually through[\[69\]](#) the use of namespaces[\[32\]](#). Generally, the resources that each environment has access to are managed by either limiting/prioritizing through Control Groups ([cgroups](#))[\[45\]](#). Resulting in an environment in which running processes, the required binaries, and their dependencies are isolated from all other running processes on the host machine and its file system. This isolated environment is referred to as a container. [\[35, 65\]](#) By sharing the same host [OS](#) kernel and containers

holding only the required binaries and libraries to run, they are considered lightweight compared to VMs because they are likely to utilize less memory and disk space. Allowing running more containers than VMs on the same physical host machine. [26, 31]

The following existing container-based virtualization tools can be used to turn applications into containers but are not limited to LXC[6], Singularity[7], and Docker[5]. For this work, the tool Docker is used to bundle the SUL into a container because it is the industry standard to date and is accompanied by a large community.[21] Docker uses so-called images (created through Dockerfiles) containing step-by-step instructions to create an image instance - namely, a container.

Thus, container-based virtualization technology is the foundation to enable scale-able SULs on-demand based on the workload at any time. Accordingly, moving the active automata learning process to a clustered system is feasible. Still, it is essential to note that the varying simultaneously running number of SUL containers must be managed.

2.4 Container Orchestration

Container-based virtualization tools provide the ability to bundle and run applications as self-contained, lightweight containers. Still, they do not offer capabilities to manage a high number of containers on a clustered system. Manual management of containers on a clustered system is challenging as the container number grows, especially across clusters with one or more physical- or virtual machines - further called nodes.

Thus, container orchestration tools provide a framework to manage containers at scale by utilizing the resources the underlying clustered system provides. Therefore, providing features to simplify the operational effort of deploying, managing the life-cycle, scaling, and resources of containers. The management of hundreds or even more containers is achievable with the help of orchestration tools. Consequently, the range of primary features regarding containers of such orchestration tools are, e.g., resource management, scheduling, fault tolerance, and auto-scaling.[33, 44, 49]

In the context of this work, the auto-scaling feature is essential to perform efficient active automata learning at scale. Auto-scaling regarding containers enables the creation and termination of containers automatically without restarting the whole system or the learning process in the context of automata learning. It is implemented via policies based on configurable thresholds. These thresholds are based on metrics such as the current utilization of CPU and memory of the containers.

Some orchestration tools allow for configuring the auto-scaling behavior or extension of the available metrics through third-party auto-scalers or custom auto-scaling policies. The benefits of auto-scaling are but are not limited to efficient resource utilization leading to reduced cost of the cluster and more available resources and meeting the demand at any time when facing a varying workload. Besides offering auto-scaling for containers, it

is in some cases possible to use this feature to scale nodes to extend the resources of the cluster, as well.

Many popular container orchestration tools can be utilized to manage containers, such as Kubernetes (K8S)[10], Docker Swarm[9], and Apache Mesos[8]. Therefore, to manage the architecture of the clustered system, the orchestration tool K8S is used due to its maturity and stability and provides the most features.[64] Due to being open-source and enabling customization, it is actively developed by a large community.[21] K8S manages the clustered system (further called K8S cluster) based on the concept of maintaining the desired state, which is described through applying, e.g., YAML files, defining K8S resources such as deployments.

To conclude, orchestrating lightweight SUL containers using K8S enables to move active automata learning to an environment of a clustered system. Furthermore, it enables the efficient use of resources and matching the number of active SULs to the varying workload of a running learning process through auto-scaling. A auto-scaling is proposed in section 3.4.

Chapter 3

Automata Learning of Clustered Systems

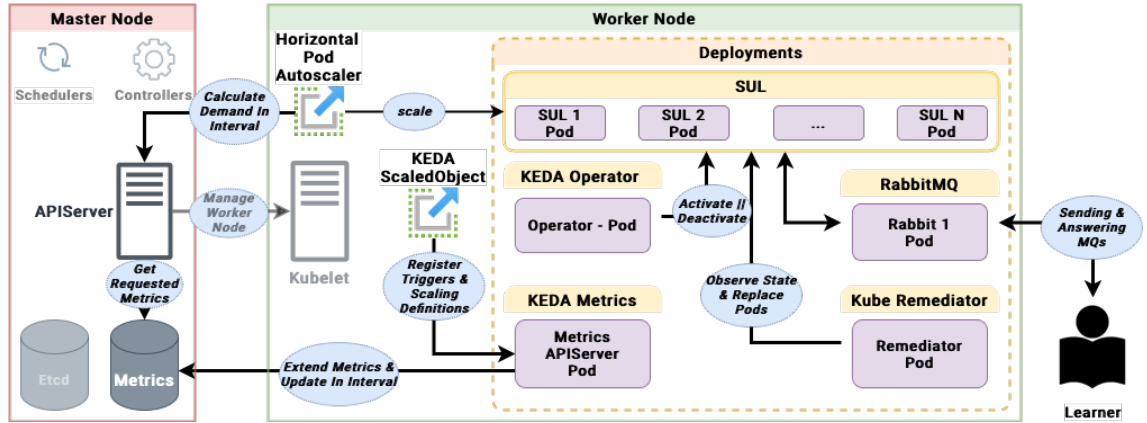


Figure 3.1: Simplified Overview Of the Architecture

The proposed architecture of this work, on which scale-able active automata learning processes are run, is illustrated in Figure 3.1. As stated in the previous chapter, the underlying cluster is managed by K8S and uses Docker as the container engine to manage and run its containers. Therefore, this architecture requires the SUL to be containerized (as a Docker image) for learning a behavioral model of the SUL. However, further consideration must be made before having a feasible solution to perform scale-able active automata learning, especially regarding real-life applications.

3.1 Membership Query Distribution

The usage of self-contained SUL containers makes direct communication between the learner and the SULs impossible. Furthermore, due to auto-scaling, the number of active SULs varies at any time. Thus, it requires a mediator who knows about each SUL to




Figure 3.2: Learner and SUL communication with RabbitMQ

handle the communication - exchanging MQs - between the two partakers. Moreover, it is required to handle communication in such a way that it is ensured that every SUL has a turn - namely, load balancing. In addition, the MQ distribution should be performed asynchronously to reduce additional overhead, e.g., wait time until a SUL receives a query. Hence a message broker is used to close the communication gap.

A message broker is a means to connect and scale a modular architecture in a distributed system by allowing communication between the parts of the architecture in question - in this case, the SULs and the learner - via so-called "message queues." Moreover, a message broker offers a common infrastructure for systems to produce and consume messages. Usually, it is based on the publish/subscribe paradigm.[54] In the context of automata learning, it is critical that such a message broker is reliable, stable, and does not introduce a noticeable overhead while exchanging MQs from the learner to the SULs and back. For this reason, the message broker RabbitMQ[22] is a suitable choice for fulfilling the stated requirements.[43]

RabbitMQ is an efficient and scale-able open-source message broker acting as a communication platform between independent applications using an Erlang-based Advanced Message Queuing Protocol (AMQP)[3], enabling asynchronous processing of data. The RabbitMQ infrastructure is built around messages, message queues, exchanges, and route bindings.[36]

The core of RabbitMQ is the concept of a message, which may contain any data. An application that sends a message through RabbitMQ is called a publisher, whereas the message receiver is called a consumer. An application may be a publisher as well as a consumer. A message always goes through an exchange, acting as a router, accepting and sending the message to the correct message queue. The exchange knows the correct

message queue based on the routing key contained in the message and the route binding connecting the message queue to the exchange. Thus, messages are always accompanied by a routing key. Applications subscribed to a message queue containing messages can consume and act upon them.

Since communication during the learning process between the learner and the **SULs** goes through RabbitMQ, it is essential to cover the existing RabbitMQ communication infrastructure used (illustrated in **Figure 3.2**). Despite the ability to create complex topologies with RabbitMQ, the infrastructure utilized is as simple as RabbitMQ can get. At its core is a direct exchange ('SUL-Direct-Exchange') through which **MQs** are passed to message queues. On the one hand, the learner publishes unprocessed **MQs**, accompanied by the routing key 'sul.input', which are distributed to the message queue 'SUL-Input-Queue,' from which **SULs** consume **MQs** for processing. On the other hand, **SULs** publish processed **MQs**, holding the routing key 'sul.output', to the message queue 'SUL-Output-Queue,' from which the learner consumes and constructs its hypothesis of the **SUL**. Consequently, the learner and **SUL** act as publishers and consumers of messages at the same time.

Input: queries \leftarrow list of unprocessed queries, *sentQueries* \leftarrow empty hashmap

Output: queries

```

1: for all query  $\in$  queries do
2:   uid  $\leftarrow$  generate random unique identifier
3:   mq  $\leftarrow$  new MembershipQuery(uid, query)
4:   add pair (uid, query) to sentQueries
5:   publish mq to 'SUL - Direct - Exchange' and routing key 'sul.input'
6: end for
7: while sentQueries is not empty do
8:   response  $\leftarrow$  consume message from 'SUL - Output - Queue'
9:   query  $\leftarrow$  get and remove query with key response.uid from sentQueries
10:  query.output  $\leftarrow$  response.output
11: end while

```

Algorithm 3.1: Simplified RabbitMQ Membership Oracle

Even though the infrastructure is simple, further adjustments must be made. A custom **MO** must be implemented and utilized by the learner that can produce and consume **MQs** RabbitMQ messages, as shown in **3.1**. Further considerations must be made regarding the custom **MO** due to the asynchronous nature of RabbitMQ. It is not guaranteed that the same order in which the learner publishes the **MQs** is adhered to when the learner consumes responses to the published **MQs**.

As a consequence, the structure of the **MQ** is marked by a Unique Identifier (**UID**) which is used to match a consumed response from any **SUL** to the appropriate **MQ**. As

shown in 3.1, each query in a batch about to be published is extended by a generated **UID**, thus making each query sent unique and, therefore, enabling the assignment of each response to the correct query.

3.2 Challenges Of Real-Life Applications

Generally, performing active automata learning to learn a target system, especially regarding real-life applications, implicate challenges. In section 2.1, few of these were mentioned. One challenge, in particular, the requirement to guarantee each query's independence, is tackled uniquely in the context of this research.

The first option is to execute a reset procedure on the target system. This functionality is commonly unavailable in real-life systems. Secondly, the only requirement is that no query affects the other, which is possible through an abstraction, e.g., using independent sessions in web applications. Thus, no direct reset functionality is required. [46]

However, a direct reset or one through abstraction may sometimes be impossible. Thus, this work performs all experiments with a type of **SUL** that terminates itself after processing one query. As a result of **K8S** trying to maintain the desired cluster state, the terminated **SUL** is restarted. Therefore, for each published **MQ**, a **SUL** container is created to process it. This type of **SUL** is further called 'disposable'-**SUL**.

Nevertheless, this approach induces other challenges while performing a learning process regarding auto-scaling and the way **K8S** tries to maintain the desired cluster state. These challenges are in-depth explained in section 3.3 and section 3.4.

3.3 Orchestrating Containerized SULs

Containers are not run directly on the cluster itself because **K8S** encapsulates containers in its basic unit called a pod. A pod can contain one or multiple containers. In this work, a pod is limited to one container; therefore, a pod is the same as one **SUL** container. In addition, a pod includes storage resources, a network IP, and other configurations on how the container(s) are run. The requested and limited computation resources (**CPU** and memory) can be set upon creation.

As mentioned above, scheduling is one of the main tasks a master node has to do. The default scheduler of **K8S** uses the information (e.g., storage and **CPU** limit) about each pod to make decisions about pod placement to ensure that not a single node in the **K8S** cluster exceeds its capacity of resources. As a result, pods surpassing their storage limit are terminated, but in contrast, surpassing the **CPU** for a short time does not lead to immediate termination of the pod.

The scheduler is challenged through the unique circumstances created by using disposable **SULs**. Replacing many **SULs** simultaneously is to be expected. Thus the scheduler is heavily burdened and may lead to an additional overhead or unexpected behavior.

Furthermore, with the scheduler in mind, it is essential to know how **K8S** manages the container life cycle. Thus, to understand the challenges introduced by using disposable **SULs** (established in [section 3.2](#)) and auto-scaling (discussed in [section 3.4](#)). Each container in a pod can be in one of the following states:

- **Waiting**: Still executes instructions to complete the container start-up
- **Running**: The container runs without problems.
- **Terminated**: The container ran to completion, failed, or is forced to terminate

In addition, it is possible to run instructions before entering a specific state through life cycle hooks. For example, the 'preStop' hook can be configured and run before a container enters the terminated state.

To conclude the life cycle management of containers, it is essential to note that each container adheres to a restart policy with limited options. The restart policy defines when to restart the containers in the pod. It becomes effective when the pod fails, succeeds, or is in an undefined state and therefore exits. The pod 'restartPolicy' field is configurable and can accept one of the following values: 'Always,' 'OnFailure,' or 'Never' (default: 'Always'). If the 'restartPolicy' becomes effective, containers enter a restart loop with an exponential back-off delay which is capped at five minutes. This delay is reset once a container runs for 10 minutes without issues. [\[14\]](#)

For managing instances of pods, **K8S** offers different resources such as deployments and jobs. The proposed architecture uses deployments to manage the **SUL** pods while learning. Deployments define long-running tasks supported by the **K8S** auto-scale feature. Furthermore, it creates a replica set that, in turn, creates the desired initial number of replicated pods (namely replicas). An image can be set, which is used to create the container inside each replica. The deployment maintains the desired number of pods specified or adjusted through auto-scaling. In addition, CPU and memory requests and limits for each pod can be set in a deployment.

It is important to note that a pod in a deployment supports only the restart policy with the set value 'Always' - this value is not configurable. This default policy results in pods continuously restarting and entering a restart loop when the pod fails or runs to completion. By the design of a **K8S** deployment, the restart loop ensures that a pod is available at any time. However, the exponential back-off delay, which is also not configurable - leads to a significant downtime between answering queries using disposable **SULs** that terminate themselves after processing one query. Consequently, each **SUL** constantly enters the restart loop after processing a query. However, this mechanism does not replace the failed

pod. Instead, the container(s) within the pod is restarted, so no heavily weighted processes are induced by this process. Still, the exponential back-off delay is a challenge to overcome.

K8S does not offer an alternative to a deployment for this work’s unique circumstances, allowing for auto-scaling. Using disposable **SULs** is still achievable while utilizing the benefits a **K8S** deployment offers. The following options may achieve this:

1. Implementing a custom **K8S** control plane enabling, e.g., the configuration of the restart policy in deployments and replacing the default **K8S** control plane
2. Instead of using **K8S** resources like deployments to orchestrate containers, implementing a custom pod management logic deployed as a pod that handles creation, termination, and scaling through the use of the **K8S** API
3. Observing each pod’s state in deployment and replacing the pods upon reaching a specific state

For this work, the decision lies on the third option because such an application already exists. For this purpose, the application Kube Remediator[15] is used to enable the use of disposable **SULs** by reducing the downtime significantly by replacing each pod upon reaching its restart loop. Although this approach does enable the use of disposable **SULs**, the act of replacing pods includes the termination of an old and the creation of a new pod. Therefore, this process may lead to an overhead.

In conclusion, **K8S** offers many features and options to orchestrate containerized applications designed for long and short-running tasks. This work holds special requirements which do not fit the **K8S** design. However, through third-party tools, it is achievable to work around the introduced challenges and make it practicable to effectively perform active automata learning on a **K8S** cluster with disposable **SULs**.

3.4 Strategies for Scaling SULs

In a scenario in which auto-scaling is required, **K8S** offers two resources to auto-scale pods: Horizontal Pod Autoscaler (**HPA**) and Vertical Pod Autoscaler (**VPA**). Even though both auto-scaler types do their task without any manual intervention, therefore, doing the task automatically, they differentiate in the way on which entity the scaling is done. The **VPA** scales the proper computation resources each pod can access based on a resource analysis over time. In contrast to the **VPA**, the **HPA** adjusts the desired quantity of replicas based on the utilization of their assigned computation resources. [29] In the context of this work, the **VPA** does not result in any significant performance gain due to the low resource requirements the used **SULs** have as a result of simply representing a Mealy machine. Because the actual goal is to process as many **MQs** as possible simultaneously, therefore, scaling up and down the number of **SUL** pods is required - thus using a **HPA**.

3.4.1 Horizontal Pod Autoscaler

The **HPA** scales a target **K8S** resource, e.g., deployment, by adjusting the desired replica count in the target resource; therefore, the **HPA** is not working directly with the pod replicas. In the example of a deployment, the maintenance of the replica count is done by the deployments replica controller, as shown in **Figure 3.3**. By default, only two basic computation resources are available as a metric to auto-scale the replica count: **CPU** and memory. [48] **HPA** scales based on the average utilization, across all replicas, of the chosen metric and triggers a replica count update through the following equation[11]:

$$desiredReplicas = \lceil currentReplicas * \frac{currentMetricValue}{desiredMetricValue} \rceil$$

Furthermore, using standard **HPA**, it is impossible to scale the replica count to zero, disabling the deployment to save resources. A simplified overview of the default **K8S** auto-scaling process, including where the aggregated metrics are pulled from, is illustrated in **Figure 3.3**.

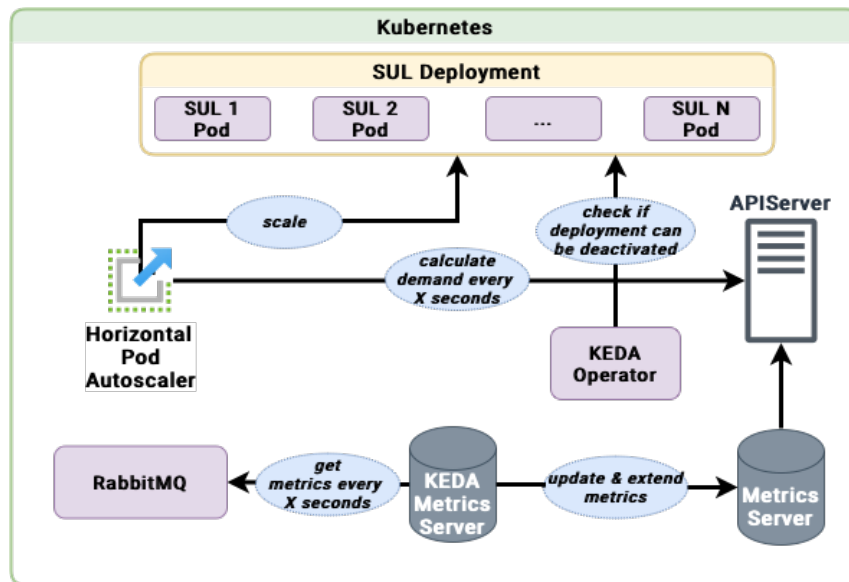


Figure 3.3: Simplified overview of the KEDA auto-scaling process.

HPA updates its metric values every 15 seconds (default value). These metrics are aggregated through the metrics server, which retrieves its information from every node in the **K8S** cluster. The aggregation is done every 60 seconds by default. Consequently, with the default settings in mind, there may be a time gap of 45 seconds in the worst case between identifying that the target should be scaled up and the scaling itself. As a result, **HPA** can miss workload peaks in a time range of seconds that might occur now and then. Furthermore, the calculation is based on the average utilization of the target metric aggregated over all pods resulting in missing out on workload peaks that occur every second. [55]

However, as mentioned above, **HPA** only offers, by default, two basic pod metrics to calculate the current demand for pod replicas. Both options are not suitable for guaranteeing the availability of a **SUL** for each **MQ** sent by the learner. Consequently, each **SUL** replica only processes one **MQ** at a time resulting in steady resource consumption for each replica. A more appropriate auto-scaling strategy is for each **MQ** in a specific RabbitMQ queue; one **SUL** replica is spawned to process that **MQ**. Consequently, the available metrics to the **HPA** must be extended by RabbitMQ-specific metrics.

As a result, a third-party solution, namely Kubernetes Event-driven Autoscaling (**KEDA**)[\[16\]](#), is utilized to enable auto-scaling for the chosen event-driven architecture: RabbitMQ.

3.4.2 Kubernetes Event-driven Autoscaling

KEDA is an extensible, lightweight component. It can be attached to a **K8S** cluster without disturbing or replacing **K8S** own auto-scaling resources like the **HPA**, therefore, being a secure option to use on any **K8S** cluster by coexisting with other auto-scaling options. It acts more like an extension to the standard **HPA** and works with it under the hood.

KEDA offers a multitude of auto-scalers regarding event-driven infrastructures using message brokers (e.g., RabbitMQ) or similar tools. Therefore allowing auto-scaling of any **K8S** resource based on the exposed metrics of each event-driven infrastructure. In the case of RabbitMQ, the following metrics (are called triggers in **KEDA**) are available for auto-scaling:

- Message Rate: Triggers auto-scaling based on the published messages to a specific queue in a second (publish rate)
- Queue Length: Triggers auto-scaling on a set number of messages in the specified queue

A **KEDA** configuration file is shown in [3.4](#) (more fields and details in [\[16\]](#)). This configuration file is used in the learning setups ([section 4.5](#)) and illustrates how a **KEDA** ScaledObject and its available fields can be configured.

Moreover, a ScaledObject offers the modification of the **HPA** behavior through the use of advanced fields in the configuration file. By default, scaling down and up can occur every 15 seconds. However, it is possible to independently change the interval for scaling up and down. Moreover, there is a difference in the way scaling up and down works. Scaling down considers a stabilization window (default: 300 seconds) in which the desired state is inferred by choosing the highest value in a list of previously calculated desired states in the set interval, e.g., in the last 5 minutes.

This window prevents a heavy alternating of the replica count and especially removing pods to create the same number of pods in the next moment. In contrast to that, scaling

Figure 3.4: KEDA ScaledObject Configuration File

```

1 spec:
2   scaleTargetRef:
3     apiVersion: apps/v1
4     # K8S resource type of the target to scale
5     kind: Deployment
6     # The name refers to the target K8S resource
7     name: sul
8   # Interval to update metrics
9   pollingInterval: 1
10  # Min. replica count to be maintained
11  minReplicaCount: 0
12  # List of auto-scale triggers
13  triggers:
14    # Type of auto-scaler
15    - type: rabbitmq
16      metadata:
17        # Protocol used to scrape the metrics
18        protocol: amqp
19        # Possible metrics: QueueLength or MessageRate
20        mode: QueueLength
21        # Threshold on which to trigger
22        # Message count or Publish/sec. rate
23        value: "1"
24        # Target RabbitMQ message queue
25        queueName: sul_input_q

```

up has, by default, no stabilization window.^[12] At last, limiting the number of pods terminated or created through scaling is also configurable.

Auto-scaling may only work for some learning setups, especially regarding scaling down replicas. It has to be ensured that a created **SUL** which consumed a **MQ**, is terminated after it finished processing the query. The following options are available:

1. Re-queueing the **MQ** before terminating the **SUL**
2. Removing the **MQ** from the queue only if it got processed by a **SUL**

The second approach is taken in this work. While this approach prevents the loss of **MQs**, it does not solve the problem that the **SUL** is terminated before processing the **MQ**. As a result, scenarios can occur in which **MQs** are passed from **SUL** to **SUL**, hindering the learning process.

K8S offers features configurable in the deployment file, solving this issue. As mentioned above, it is possible to configure a 'preStop' life cycle hook. It runs before the container is terminated. Through this hook, it is possible to delay the termination until the **SUL** in question completes processing the **MQ**. However, this approach requires that the container offers a shell environment and eventually read/write access (based on the implementation of the hook) on the **K8S** cluster it runs on. The native image created with GraalVM offers, by default, no shell. Therefore, this approach is not suitable.

In this work, the following approach is taken. Pods about to be terminated are going through a grace period, allowing for a graceful termination. Meaning **K8S** allows the pod to shut down, e.g., RabbitMQ connection, or let it finish processing and terminate itself before being forcefully killed by **K8S**. The grace period is, by default, 30 seconds. However,

the grace period has to be adjusted while using applications that run heavy processing tasks that exceed this boundary.

To summarize, auto-scaling active automata learning with the design choices made in this work is achievable through the dedicated event-driven auto scaler **KEDA**. Therefore, having one **SUL** replica ready for each **MQ** in a specific RabbitMQ queue may lead to more efficient use of computational resources than the approach of using a constant number of **SULs**. In contrast, the auto-scaling strategy might be slower regarding the total execution time. Furthermore, auto-scaling implicates additional management effort regarding the container life cycle depending on the learn setup.

Chapter 4

Evaluation Method

The architecture explained in [chapter 3](#) and shown in [Figure 3.1](#) has been set up and is used to perform the active automata learning setups introduced in this chapter. The architecture is constructed as a single node cluster with the help of MicroK8s[\[17\]](#) and is equipped with the following hardware - 126 GB of RAM and an AMD EPYC™ 7443P [CPU](#) with 24 cores @ 2.85 GHZ with 48 threads (more details on [\[1\]](#)).

4.1 Experiment: Simple Coffee Machine

Evaluation of the proposed architecture regarding practicability, performance, auto-scaling, and limitations is done by executing active automata learning processes. In [section 2.2](#), the Mealy machine representing a simple coffee machine, was introduced. This Mealy machine is containerized and used as a disposable [SUL](#) in the different learning setups to evaluate the architecture.

Moreover, as mentioned in [subsection 2.1.1](#), LearnLib offers various algorithms and filters regarding the learning process. The learner used in the learning setups runs the learning process with the following LearnLib components:

- Query cache that reduces the total sent queries by filtering out queries that were already sent
- The Equivalence Oracle simulates the behavioral model of the [SUL](#); thus, no approximation of [EQs](#) takes place, reducing the total amount of sent queries significantly.
- LearnLib offers many learning algorithms. In this work, the Direct Hypothesis Construction ([DHC](#)) learning algorithm is used because it supports query batches, enabling query distribution for simultaneous processing. Moreover, the primary reason is that the [DHC](#) algorithm favors big batch sizes compared to the other learning algorithms available in LearnLib, such as L^* .[\[46\]](#)

As a result, the following metadata of the Mealy coffee machine in use is illustrated in Table 4.1.

<u>Data</u>	<u>Value</u>
<i>Number Of States</i>	<i>5</i>
<i>Number Of Inputs</i>	<i>4</i>
<i>Number Of Outputs</i>	<i>3</i>
<i>Number Of Queries Sent</i>	<i>118</i>
<i>Min. Batch Length</i>	<i>1</i>
<i>Max. Batch Length</i>	<i>4</i>
<i>Avg. Batch Length</i>	<i>1</i>
<i>Min. Sequence Length</i>	<i>1</i>
<i>Max. Sequence Length</i>	<i>6</i>
<i>Avg. Sequence Length</i>	<i>3</i>

Table 4.1: Coffee Mealy Machine Statistics

4.2 SUL Containerization

To orchestrate **SUL** instances with **K8S**, turning the **SUL** application into a container on which the experiments are based is required. Although creating a container is not complicated when done with a tool like Docker, the generated overhead in the process can vary significantly based on the binary and dependencies required to run the **SUL**. Reducing the overhead is vital in this work because many containers are expected to be terminated and new ones created over time when using auto-scaling and disposable **SULs**. Consequently, the underlying foundation of the **SUL** container determines the impact on:

- Processing time of a **MQ**
- Scale-ability of the **SUL**, including the time it takes to create, start up, and terminate a container
- Deployment time can vary by the size of the container

Thus it can have quite an impact on the total execution time and practicability of a learn setup. Therefore, crucial details about the implementation of the **SUL** representing the simple coffee machine are as follows:

- LearnLib is used to implement the coffee machine’s behavioral model represented as a Mealy machine
- The language that is used is Java[13]. As a result, the **SUL** runs by default on a Java Virtual Machine (**JVM**)[67]

- Through the use of the framework Spring Boot[23] or Spring Native[4] (explained below) to trivialize the creation process of an application that is runnable, and the use of third-party libraries
- The RabbitMQ third-party library for the Spring Boot framework

These details may notably impact image size, application start-up time, and, therefore, the total execution time of a learning process.

The Docker container layer wrapped around the application itself adds a not noticeable overhead of approximately 0.036s on the container creation time, as stated by the author of [38].

The size of the Docker image is a concern when using Docker image registries, where the image is pulled from to create containers. A smaller image size results in a faster load speed from the registries in question and thus leads to faster container creation. Table 4.2 illustrates the variance a coffee machine SUL image size can have used different base images required to run the coffee machine and its dependencies. At last, the start-up time of the

<u>Base Image</u>	<u>Compressed Size</u>
<i>openjdk:11</i>	<i>334,91 MB</i>
<i>openjdk:11-jdk-slim</i>	<i>241,63 MB</i>
<i>adoptopenjdk</i> <i>/openjdk11:alpine-jre</i>	<i>67,74 MB</i>
<i>native-image graalvm</i>	<i>26,02 MB</i>

Table 4.2: Image Size Variance: Example Coffee Machine

SUL itself determines the time it takes until the SUL is ready to accept MQs. As shown in Table 4.3, the time it takes to start up the coffee machine SUL itself can fluctuate and vary a lot based on the implementation, how the executable is built and how the Docker image is created. In this work, a significant reduction of the start-up time and image size could

<u>Base Image</u>	<u>Average Start Up Time</u>
<i>adoptopenjdk</i> <i>/openjdk11:alpine-jre</i>	3.1796s
<i>native-image (GraalVM)</i>	0.0982s

Table 4.3: Start-Up Time Coffee Machine SUL - JVM vs. GraalVM

be achieved by building a native image with GraalVM. Therefore, reducing the overall overhead induced by containerization and of the SUL application itself. Thus, the native image is used to create coffee machine SULs in a learning process.

4.3 Simulating Real-Life Systems

The processing time of a **MQ** regarding the Mealy machine implementation is near instant, in contrast to real-life applications, which may take noticeably longer, even seconds. Consequently, adjustments are made in the learning processes to accommodate real-life systems, although it does not replace them.

The accommodation to real-life systems is an artificial delay to each **SUL** processing time. The artificial delay is calculated based on the sequence length of the **MQ** to be processed. For each symbol in the sequence, a random time interval is accumulated, resulting in the wait time before the **MQ** is processed. The time interval used in the calculation is configurable, and various values are used in this work. See [section 4.5](#) for more details on the values used. This approach simulates real-life applications designed to process heavy tasks that take a long time.

As introduced in [section 3.2](#), using the disposable **SULs** accommodates real-life systems that do not offer functionality or take a long time to reset, fulfilling the requirement of **MQ** independence.

Both approaches merged cover complex and extensive systems that may include many services and persistent storage solutions that take longer to process queries and often do not offer functionality to reset the whole system. These systems are commonly built so that a reboot is needed to set the system to its initial state because they are not primarily built for active automata learning.

4.4 Data Acquisition

In this work, for each execution of a learn setup (explained in [section 4.5](#)), data about the **SUL** and the learning process is collected. Besides the already mentioned details about the coffee machine mentioned in [Table 4.1](#), additional data is acquired for evaluation:

- Learn setup used
- Processing time (min, max, avg) of each **SUL**
- Start-up time (min, max, avg) of each **SUL**
- Interval between each response the learner gets (min, max, avg)
- Time it takes to process a batch of **MQs** (min, max, avg)
- Artificial delay accumulated (min, max, avg) over each input in a **MQ** sequence of each **SUL** before the query is processed
- The total execution time it took to learn the model

Furthermore, the learner calculates and accumulates the data and receives **SUL** specific data as an appendix to the **MQ** response. At the end of each learning process, the data is averaged. Additionally, it is ensured that the executed learning process sends the same queries in the same order for each learning setup.

4.5 Learning Setups

The constructed architecture allows configuring of various parameters of a learning setup, such as:

1. Configuration of the **SUL** deployment including:
 - The **SUL** image that should be used. Determines the behavior of the **SUL**, e.g., coffee machine
 - Minimum and maximum **SUL** instances, which may be overwritten by **KEDA** when using auto-scaling
 - Deployment-related configuration, which may be unique to the **SUL** requirements, e.g., **termination grace period** or a pre-stop hook
2. If the deployment is to scale automatically, a **KEDA** ScaledObject must be configured including:
 - **SUL** instance count range limit that the auto-scaling does not exceed or undercut
 - The event-driven infrastructure used. In this case, RabbitMQ is used, and the auto-scale triggers:
 - Metrics: 'QueueLength' or 'MessageRate'
 - Target queue
 - Trigger threshold value determining when to auto-scale
 - Polling interval of the metrics and cool-down period for disabling the deployment when no new message arrived in the target queue
 - For more fine-grained control over the auto-scaling process, the **HPA** can be configured based on the requirements of the learning setup

In this work, the learning setups can be differentiated into two categories. The first category of learning setups uses a constant number of disposable **SULs**. The following **SUL** counts are used: 1, 2, 4, 8, 16, 32, and 64. The purpose of this category is to highlight the scalability and practicability of the approach taken.

The second category uses auto-scaling with **KEDA** based on the queue length of the **SUL** input queue. For each **MQ** in the queue at the time, the new demand is calculated

by the **HPA**. In this case, **KEDA** is configured, so the metrics are updated each second. The following **HPA** demand calculation intervals in seconds are used: 5, 15, and 30. Additionally, one demand calculation approach is taken in which the demand of scaling up is calculated every five seconds and for scaling down every 30 seconds. The different **HPA** intervals may impact the total learning time. Furthermore, this category is run with zero and one initial **SULs**.

Before the learning process of a learning setup is executed by sending the first batch of queries, it is ensured that the mentioned initial number of **SULs** are deployed and are ready to receive queries.

The termination grace period for each learning setup is set to 10 seconds, and for more extensive delays, 60 seconds. These values are sufficient for the learning setups in this work, guaranteeing that each **MQ** is processed by the **SUL** before it is forced to terminate through **K8S**.

As mentioned in [section 4.3](#), an artificial delay is added to the processing time of each **SUL**. A random number in an interval is calculated for each input in a sequence of a **MQ**. Furthermore, the calculated numbers are accumulated, resulting in the processing delay. The following millisecond intervals are used: 100-150, 500-1000, 3000-5000, and 5000-10000.

Finally, the evaluation baseline used in [chapter 5](#) for comparison is the learning setup with one disposable **SUL**. Furthermore, both categories mentioned above are compared to each other regarding total execution time and resources spent.

Chapter 5

Results

In this section only the summary of the results are shown. The complete overview of results, including all details, can be found in [Appendix A](#).

5.1 Architecture Overhead

<u>Overhead In Seconds</u>	<u>1 SUL</u>	<u>8 SULS</u>
<u>Minimum</u>	0,007742889667	0,083867635
<u>Average</u>	9,949262212	1,737595732
<u>Maximum</u>	24,18339815	12,61747477

Table 5.1: Overhead In Seconds Until SUL Responds

Before diving into the performance gains of the different strategies, it is essential to note the overall overhead that the architecture induces. [Table 5.1](#) demonstrates the minimum, average, and maximum overhead in seconds that is added to the time it takes until a [SUL](#) responds to a query. Observing a single [SUL](#) for the whole learning process leads to the following findings:

1. The minimum overhead is in the nanoseconds ranges similar to the start-up time of the coffee machine application itself, as shown in [Table 4.1](#). The reason is the behavior of the restart loop enforced by [K8S](#); the first restart occurs immediately. The restart loop does not terminate the pod in question, so there is no need to create a new pod. Instead, the container in the pod is restarted, and therefore the coffee machine application itself. Leading to a quick start-up and no overhead at all.
2. Overhead is introduced since containers need to be terminated and restarted. As mentioned in [section 3.3](#), the application Kube Remediator is used to terminate pods that reach the first stage of the exponential back-off in the restart loop. As a result, the [K8S](#) deployment maintained the desired [SUL](#) count by restarting a new




Figure 5.1: % Change In Start-Up Time Based On SUL Count - Low Delay

pod in place of the terminated one. Consequently, the termination and creation of a pod introduce an average overhead of approximately 10 seconds until a **SUL** is ready to receive a query.

3. The maximum overhead close to 25 seconds is because **K8S** sometimes fails to create a pod for various reasons. Therefore, the pod is terminated and created again, leading to considerable overhead.

However, the introduced overhead can be mitigated using more, e.g., eight **SULs** as shown in [Table 5.1](#). It is important to note that more **SULs** do not reduce the overhead of the architecture affecting each **SUL**. Instead, the increasing count reduces the time interval until a **SUL** is ready to respond to a query.

Another finding is that the more **SULs** are used in a learning process, the overhead of **K8S** itself increases. This overhead is more prevalent in using a learning setup with low to no processing delay ([Figure 5.1](#)) compared to one with a high processing delay ([Figure 5.2](#)). The reason is that the management effort to schedule, reserve resources, create and terminate pods is much higher in low-delay learning setups due to the frequency with which pods are replaced. Leading to an average slower start-up time of approximately 45% at 64 **SULs**. In this case, the **K8S** cluster starts to throttle the resources to which the **SULs** have access. Thus, having more resources available to maintain the desired cluster state. The big outliers while using 32 and 64 **SULs** in both graphs maybe since **K8S** uses

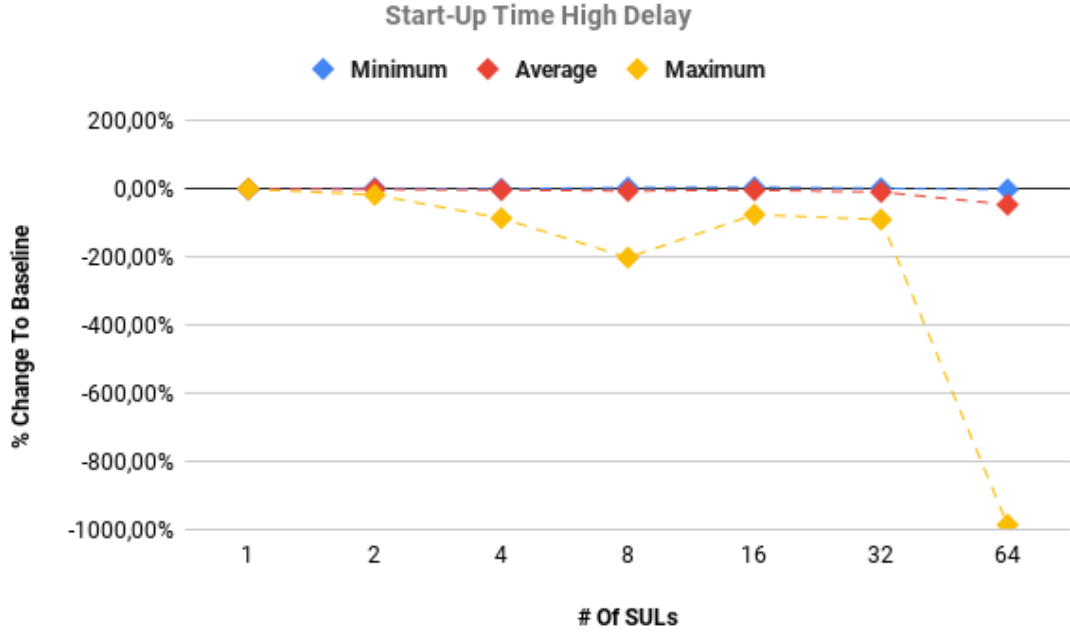


Figure 5.2: % Change In Start-Up Time Based On SUL Count - High Delay

many resources to maintain, schedule, terminate, and create many **SULs** at once. The maximum overhead fluctuates a lot.

5.2 Strategy: Constant Number Of SULs

Besides the introduced overhead through the proposed architecture, while using disposable **SULs**, parallelization is still very effective in improving the total learning time, as illustrated in [Figure 5.3](#). Using, e.g., 64 **SUL** instances increase the performance gain by a factor of two compared to the baseline of one **SUL**. Even using only two **SULs** leads to a 50% performance increase. Therefore, the results suggest that learning target systems that respond fast to queries benefit significantly by an increasing **SUL** count. However, the performance increase for each additional **SUL** reduces significantly. At about the point of using eight **SULs**, multiply the **SUL** count by a factor of two the performance increases by approximately 5% on average.

An assumption can be made that the maximum performance gain can be achieved by using as many **SULs** as the total queries sent to the target system over the whole learning process. Therefore, reducing the time till a **SUL** is ready to respond to zero. Consequently, a **SUL** is always ready to process each query sent. In the case of the simple coffee machine, the maximum performance gain would be at 118 **SULs**. Learning the coffee machine would be around a second using this many **SULs**. However, **K8S** has a pod limit of 110 per node, and the experiments are run on a single node cluster.

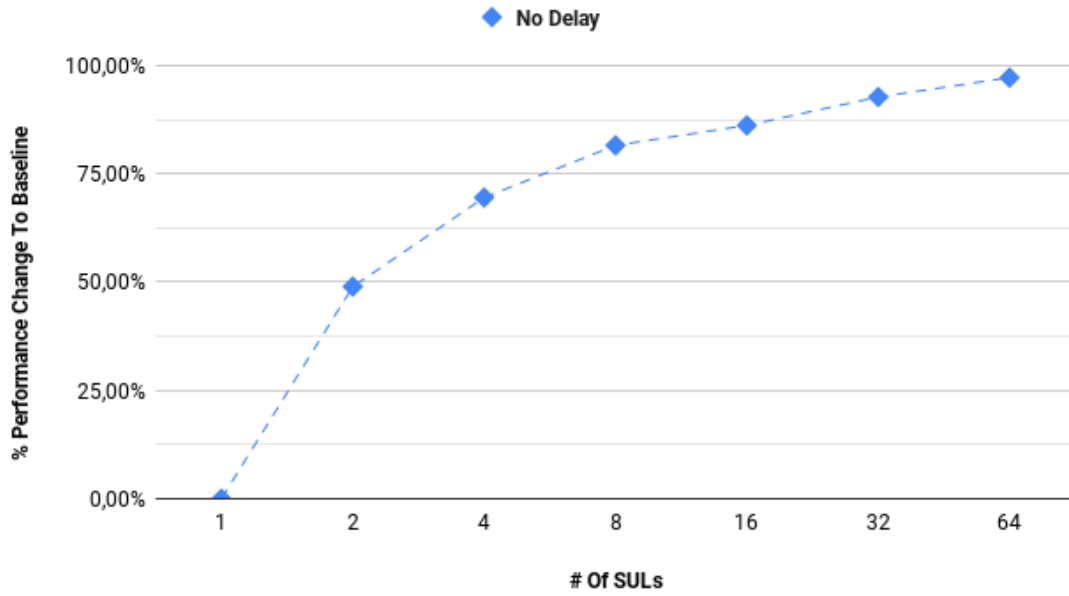


Figure 5.3: % Performance Gain - No Delay

While it would be feasible to do this with a very light and fast target system, such as the coffee machine used in this work, using a complex real-life system as the learning target would not be feasible. The reason is that increasing the **SUL** count by a factor also increases the used resources and, therefore, the cost of maintaining the cluster by the same factor.

The graph suggests that the best cost-benefit ratio would be around six to eight **SULs** active at a time leading to a performance gain of around 70-77% when learning systems that are fast to start up and fast to respond. In contrast, different results can be observed introducing processing time delay to each **SUL**.

5.2.1 Introducing Processing Time Delay

The reduction of the cost-benefit for each **SUL** is even more prevalent when using systems that respond slowly to queries - in this work, represented adding processing delays. Using sixteen **SULs** instead of 8 improves the performance by under 1% when using a delay interval such as 5000ms-10000ms, resulting in the worst case in a processing time of 60 seconds using the coffee machine with a maximum sequence length of 6 (Table 4.1). The higher the processing delay, the more negligible the performance gain of adding more **SULs**. Figure 5.4 suggests that the performance gain is limited at approximately 77%.

The reason is that the **SULs** take so long to respond to a query that fewer **SULs** are needed to fill the gap until a **SUL** is ready. Since a query's processing time takes longer than terminating and starting up a new **SUL**, using more **SULs** than eight is a waste of

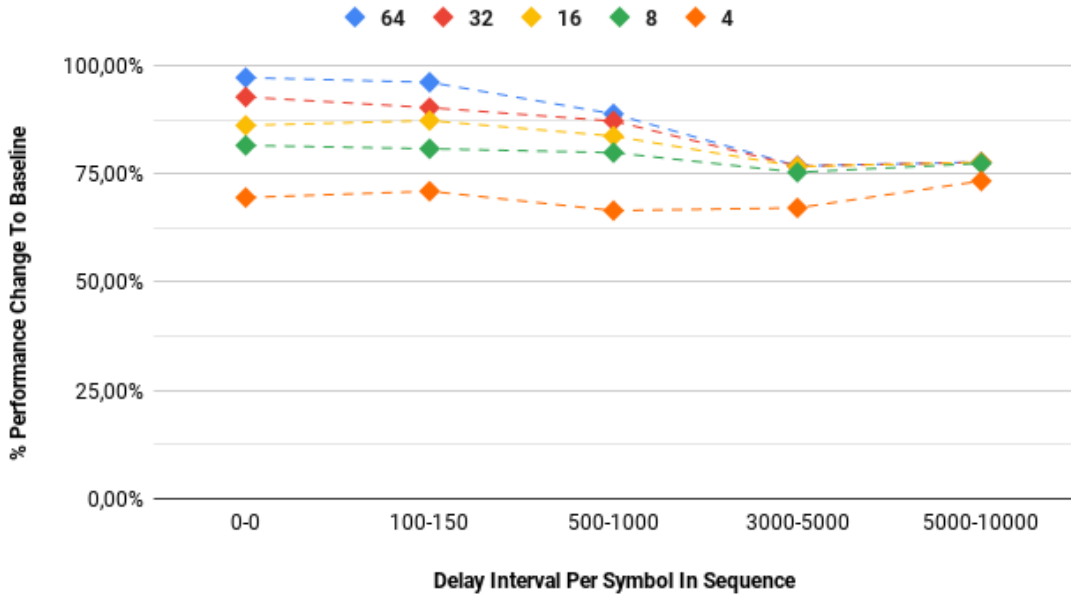


Figure 5.4: % Performance Gain - With Processing Delay

resources in the high delay area; no performance gain is being made. It gets to the point that the performance gain of four **SULs** are close to the one of 64, 32, 16, and 8 **SULs**, as emphasized in the graph.

An assumption can be made here that the performance gain is capped around the maximum query batch size of the experiment. It is four in the case of the used coffee machine, as shown in [Table 4.1](#). Nevertheless, the coffee machine has an average batch size of one. Therefore, using a constant number of **SULs** always leads to time intervals in which most **SULs** are idle, and resources are wasted.

In contrast to that, in the scenario that all **SULs** process their queries nearly at the same time (in approximately 10 seconds as shown in [Table 5.1](#)), using more **SUL** instances more the maximum batch size can be benefited from. New queries are being sent, and all available **SULs** are rebooting; therefore, the process lies dormant for some seconds. New queries can be processed with more **SULs** while the others reboot. However, this may lead to time intervals with idle **SULs**. The results suggest that the optimum cost-benefit ratio lies at around 6-8 **SUL** instances.

5.3 Strategy: Auto-Scaling

As illustrated in [Figure 5.5](#), different scaling intervals were used. The graph demonstrates a high fluctuation in performance gain across all scaling intervals compared to the baseline. The results suggest that the scaling strategy has no noticeable impact on the performance gain using the coffee machine as the target system. The outlier emphasizes the fluctuation

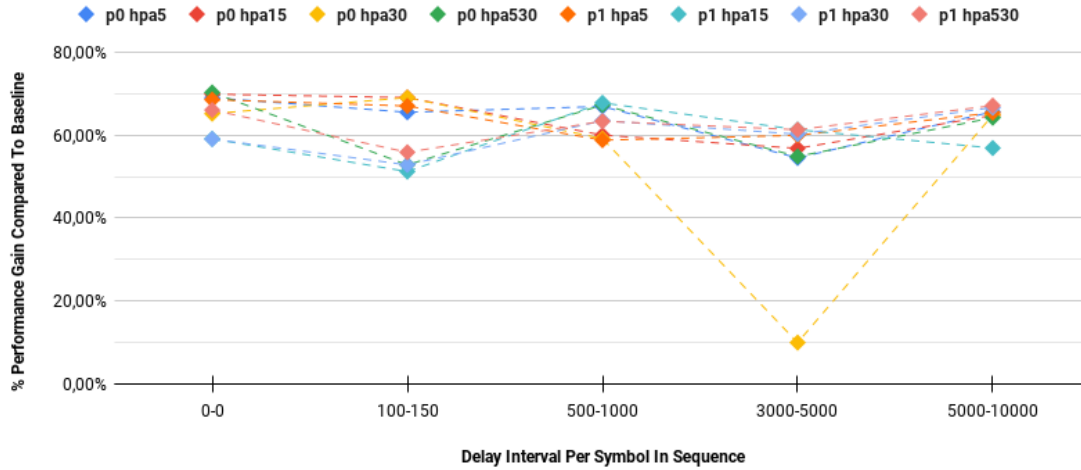


Figure 5.5: Comparison Of Scaling Intervals

because it may lead to no significant improvement in the worst case compared to using only one **SUL**. This case may be caused due to wrong auto-scaling timings and other factors such as scheduling.

Furthermore, each scaling interval is used with no and one initial **SUL** before the learning process starts, as illustrated in Figure 5.5 by the strategy prefix *p0* and *p1* for no and one initial **SUL** respectively. On average, using no initial **SUL** results in a better performance of 1-2% compared to one initial **SUL**. It results from the fact that one initial **SUL** directly removes one **MQ** from the queue before the first auto-scaling takes place. At this time, one or more messages are already processed, and **SUL** count is scaled less compared to having no initial **SULs**. The first batch containing four queries may lead to a count of one to three instead of four **SULs** after the first auto-scaling takes place. The result of this scenario depends, e.g., on the scaling interval used and processing time of the **SUL**, as emphasized by Figure 5.6.

However, the strategy of using one **SUL** for each **MQ** in the RabbitMQ at the time of the demand calculation leads to good results compared to the constant **SUL** strategy; especially in the context of high processing times (Figure 5.6). The auto-scaling performance gain is more reliable than the other strategy across all processing delay intervals. On average, a performance gain of approximately 65% to 70% could be achieved using auto-scaling while maintaining an excellent cost-benefit ratio by significantly reducing idle **SULs**.

5.4 Evaluation

Performing active automata learning on a clustered system by parallelizing and distributing the **MQs** to numerous running **SUL** instances where the **MQs** get simultaneously processed can lead to a notable reduction in the total execution time. Depending on the internal/ex-




Figure 5.6: KEDA Compared To Constant Number Of SULs

ternal factors of the actual learning setup, e.g., the complexity, time to reset, processing time, and the start-up time of the **SUL**, the performance is expected to fluctuate.

Using a constant number of **SULs** often leads to intervals in which **SULs** idle, thus, leading to wasted resources and an avoidable overhead that is induced by maintaining the idle **SULs**. These time intervals and the number of idle **SULs** are reduced significantly by using auto-scaling strategies, therefore, using the resources more efficiently and reducing the cost.

Furthermore, it depends on the constructed architecture specification and the clustered system used for the learning process. In the case of this work, a change to the architecture that only restarts the container itself would significantly reduce the overhead mentioned above to a level of only the start-up time of the application itself matters. Replacing pods induces heavily weighted tasks such as, but not limited to, scheduling, resource reservation, and creating the container itself. Thus, leading to significant overhead.

Moreover, increasing the **SUL** instance count does not always result in increased performance. Instead, it can lead to no performance gain at all. There are two reasons for this:

- First, the overhead of the used technologies grows as the container management task increases in complexity or several **SULs** are fast enough so that each addition might result in an idle **SUL** and no performance gain.

- Second is that the **MQs** are sent in batches, and each batch has to be fully processed until the next batch of **MQs** is sent. Thus there is a constraint on the performance dependent on the query batch size the learner sends. It is important to note that more extensive batch-size experiments may yield a better performance gain than the coffee-machine experiment.

Finally, the number of **SUL** instances is limited by the resource consumption of each instance and the available resources of the clustered system in question. Consequently, all those constraints must be considered when constructing a learning setup and can affect the total execution time of a learning process.

To conclude, the performance gain is limited by the batch size sent by the learner. However, an auto-scaling strategy that allows for more **SULs** than the maximum batch size would increase the performance further. In the case of this work, a strategy that scales up two **SULs** for every **MQ** would perform better and more efficiently. As the data suggests, six to eight pods would be the optimum in both high and low-delay areas.

Chapter 6

Conclusion

This work researched how active automata learning using LearnLib can be performed on a clustered system by distributing test queries to multiple **SULs** to process them simultaneously. An architecture was constructed with the tools **K8S**, Docker, and an event-driven auto-scaler named **KEDA** to implement the auto-scaling strategy. Furthermore, accommodations were made, such as processing delays and restarting the **SUL** instances, guaranteeing query independence and, t, simulating more complex real-life systems. Therefore, it elaborated extensively on two strategies for reducing the total learning time of a model, namely a constant number of **SULs** and auto-scaling. Moreover, both strategies are compared to each other regarding the efficient use of resources and performance gain.

In theory, both strategies are beneficial to reduce the overall learning time. However, using a constant number of **SULs** is inefficient regarding resource use and can not be adjusted without restarting the learning process to accommodate a varying workload. Depending on the target system, using high **SUL** count may be beneficial but, in contrast, may even lead to no benefit, only more costs. Thus, without knowing the specifications of the learning process of the target system, such as batch size and processing time, it is difficult to set the correct number of **SULs** to work with to get an excellent cost-benefit ratio. Besides, the query batch size might heavily fluctuate, leading to many idle **SULs** in the worst case. However, this strategy could achieve a performance gain near a factor of two.

Utilizing auto-scaling turned out to be the more reliable choice and resulted in a good performance for simple applications and complex systems. Besides that, a slight change to the auto-scaling strategy would result in a more optimal learning process regarding the total learning time and resource usage.

To conclude, the proposed architecture enables efficient active automata learning of different types of target systems. It is a feasible and practical foundation with limitations that enables learning of complex real-life applications if improved upon. Thus, several ideas

primarily focused on real-life systems are provided to extend and improve the proposed architecture for future research.

6.1 Future Research

Learning Real-Life Systems Enabled Through Mappers Although this work tries to get close as possible to real-life systems through using disposable **SULs** and additional delays, this approach does not replace them. The paper [59] introduces so-called "mappers" that serve as an interface for communication with real-life systems such as web applications. The tool **ALEX**[2] utilizes this concept to perform active automata learning with LearnLib on web applications. Opening up a way to configure **SUL** mappers on top of the architecture enables learning models of real-life systems. Another approach would be to apply the architecture discovered in this work to the tool **ALEX**. Extending the proposed architecture with mappers may be achieved by deploying an application that accepts each **MQ** and converts them to queries understandable by the target system before being sent. Furthermore, the mapping itself might be implementable through configuration files, namely **K8S** configuration maps.

Reducing The Architecture Overhead Besides using a slightly better auto-scaling strategy, a more significant improvement in the total learning time can be made by reducing the overhead of the proposed architecture. As mentioned in **chapter 5**, the Kube Remediator replaces pods in the restart loop with exponential back-off. Although this tool enables the use of disposable **SULs**, it introduces significant overhead, heavily slowing down the learning process. One way to achieve this is by modifying the Kubelet code to enable the configuration of the restart loop or modifying the exponential back-off itself, then compiling the custom Kubelet binary and applying it to the nodes on which the architecture is running. Thus, removing the Kube Remediator third-party dependency and the need to replace pods theoretically results in a higher performance gain through removing the introduced overhead.

Asynchronous Sending of Membership Queries and Equivalence Queries As mentioned, learning time reduction is related to or even limited by the query batch size. The reason is that the learner has to wait until a whole batch is processed and responded to until new queries are sent. However, the **DHC** algorithm leads to the largest batch size compared to the other learning algorithms available, an asynchronous algorithm or, instead, an algorithm that does not wait till all queries in a batch are processed before new queries are approximated. This approach has, in theory, the potential to increase performance further. Additionally, a learner using an asynchronous algorithm can use the message broker more effectively.

Appendix A

Index Of Abbreviations

DHC Direct Hypothesis Construction	23
SUT System Under Test	1
SUL System Under Learning	1
FSM Finite-State Machine	1
MO Membership Oracle	7
MQ Membership Query	2
MBT Model-Based Testing	1
MAT Minimally Adequate Teacher	5
EO Equivalence Oracle	7
EQ Equivalence Query	2
AMQP Advanced Message Queuing Protocol	14

VM Virtual Machine	9
OS Operating System	9
HPA Horizontal Pod Autoscaler	18
cgroups Control Groups	9
COS Container-based operating system	9
K8S Kubernetes	11
CPU Central Processing Unit	9
UID Unique Identifier	15
JVM Java Virtual Machine	24
VPA Vertical Pod Autoscaler	18
KEDA Kubernetes Event-driven Autoscaling	20
ALEX Automata Learning Experience	2
NGLL Next Generation LearnLib	2

List of Figures

2.1	Learn Process With LearnLib Based On [37]	6
2.2	A simple coffee machine as a Mealy machine [60]	8
3.1	Simplified Overview Of the Architecture	13
3.2	Learner and SUL communication with RabbitMQ	14
3.3	Simplified overview of the KEDA auto-scaling process.	19
3.4	KEDA ScaledObject Configuration File	21
5.1	% Change In Start-Up Time Based On SUL Count - Low Delay	30
5.2	% Change In Start-Up Time Based On SUL Count - High Delay	31
5.3	% Performance Gain - No Delay	32
5.4	% Performance Gain - With Processing Delay	33
5.5	Comparison Of Scaling Intervals	34
5.6	KEDA Compared To Constant Number Of SULs	35

Index Of Algorithms

3.1 Simplified RabbitMQ Membership Oracle	15
---	----

Bibliography

- [1] Amd epyc™ 7443p. <https://www.amd.com/de/products/cpu/amd-epyc-7443p>. Accessed: 2022-10-04.
- [2] Automata learning experience (alex). <https://learnlib.de/projects/alex/>, note = Accessed: 2022-09-14.
- [3] Cloudamqp official website. <https://www.cloudamqp.com/>, note = Accessed: 2022-03-31.
- [4] Compiling spring applications to native executables. <https://docs.spring.io/spring-native/docs/current/reference/htmlsingle/>, note = Accessed: 2022-09-14.
- [5] Container-based operating systems virtualization - docker. <https://www.docker.com>, note = Accessed: 2022-09-14.
- [6] Container-based operating systems virtualization - lxc. <https://linuxcontainers.org>, note = Accessed: 2022-09-14.
- [7] Container-based operating systems virtualization - singularity. <https://docs.sylabs.io/guides/3.5/user-guide/introduction.html>, note = Accessed: 2022-09-14.
- [8] Container orchestration tool - apache mesos. <https://mesos.apache.org>, note = Accessed: 2022-09-14.
- [9] Container orchestration tool - docker swarm. <https://docs.docker.com/engine/swarm/>, note = Accessed: 2022-09-14.
- [10] Container orchestration tool - kubernetes. <https://kubernetes.io>, note = Accessed: 2022-09-14.
- [11] Container orchestration tool - kubernetes - hpa. <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>, note = Accessed: 2022-09-29.

- [12] Horizontal pod autoscaling. <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/#configurable-scaling-behavior>. Accessed: 2022-10-04.
- [13] Java official website. <https://www.java.com/de/>. Accessed: 2022-10-04.
- [14] K8s pod lifecycle. <https://kubernetes.io/docs/concepts/workloads/pods/pod-lifecycle/>. Accessed: 2022-10-04.
- [15] Kube remediator. https://github.com/ankilosaurus/kube_remediator. Accessed: 2022-10-04.
- [16] Kubernetes event-driven autoscaling. <https://keda.sh>, note = Accessed: 2022-09-14.
- [17] Kubernetes setup tool - microk8s. <https://microk8s.io>, note = Accessed: 2022-09-14.
- [18] Learnlib official website. <https://learnlib.de/>. Accessed: 2022-03-31.
- [19] Learnlib official website - performance. <https://learnlib.de/performance/>. Accessed: 2022-10-04.
- [20] libalf official website. <http://libalf.informatik.rwth-aachen.de>. Accessed: 2022-10-04.
- [21] Most loved and dreaded technologies. <https://survey.stackoverflow.co/2022/#most-loved-dreaded-and-wanted-tools-tech-love-dread>. Accessed: 2022-10-04.
- [22] Rabbitmq official website. <https://www.rabbitmq.com/>. Accessed: 2021-11-21.
- [23] Spring boot official website. <https://spring.io/projects/spring-boot>. Accessed: 2022-10-04.
- [24] Fides Aarts, Joeri De Ruiter, and Erik Poll. Formal models of bank cards for free. In *2013 IEEE Sixth International Conference on Software Testing, Verification and Validation Workshops*, pages 461–468. IEEE, 2013.
- [25] Fides Aarts, Julien Schmaltz, and Frits Vaandrager. Inference and abstraction of the biometric passport. In *International Symposium On Leveraging Applications of Formal Methods, Verification and Validation*, pages 673–686. Springer, 2010.
- [26] Theodora Adufu, Jieun Choi, and Yoonhee Kim. Is container-based technology a winner for high performance scientific applications? In *2015 17th Asia-Pacific Network Operations and Management Symposium (APNOMS)*, pages 507–510. IEEE, 2015.

- [27] S Anish Babu, MJ Hareesh, John Paul Martin, Sijo Cherian, and Yedhu Sastri. System performance evaluation of para virtualization, container virtualization, and full virtualization using xen, openvz, and xenserver. In *2014 fourth international conference on advances in computing and communications*, pages 247–250. IEEE, 2014.
- [28] Naylor G Bachiega, Paulo SL Souza, Sarita M Bruschi, and Simone Do RS De Souza. Container-based performance evaluation: a survey and challenges. In *2018 IEEE International Conference on Cloud Engineering (IC2E)*, pages 398–403. IEEE, 2018.
- [29] David Balla, Csaba Simon, and Markosz Maliosz. Adaptive scaling of kubernetes pods. In *NOMS 2020-2020 IEEE/IFIP Network Operations and Management Symposium*, pages 1–5. IEEE, 2020.
- [30] Oliver Bauer, Johannes Neubauer, Bernhard Steffen, and Falk Howar. Reusing system states by active learning algorithms. In *International Workshop on Eternal Systems*, pages 61–78. Springer, 2011.
- [31] Aditya Bhardwaj and C Rama Krishna. Virtualization in cloud computing: Moving from hypervisor to containerization—a survey. *Arabian Journal for Science and Engineering*, 46(9):8585–8601, 2021.
- [32] Eric W Biederman and Linux Networx. Multiple instances of the global linux namespaces. In *Proceedings of the Linux Symposium*, volume 1, pages 101–112. Citeseer, 2006.
- [33] Emiliano Casalicchio. Container orchestration: a survey. *Systems Modeling: Methodologies and Tools*, pages 221–235, 2019.
- [34] Susanta Nanda Tzi-cker Chiueh and Stony Brook. A survey on virtualization technologies. *Rpe Report*, 142, 2005.
- [35] Minh Thanh Chung, Nguyen Quang-Hung, Manh-Thin Nguyen, and Nam Thoai. Using docker in high performance computing applications. In *2016 IEEE Sixth International Conference on Communications and Electronics (ICCE)*, pages 52–57. IEEE, 2016.
- [36] Philippe Dobbelaere and Kyumars Sheykh Esmaili. Kafka versus rabbitmq: A comparative study of two industry reference publish/subscribe implementations: Industry paper. In *Proceedings of the 11th ACM international conference on distributed and event-based systems*, pages 227–238, 2017.
- [37] Marco Henrix. *Performance improvement in automata learning*. PhD thesis, Master thesis, Radboud University, Nijmegen, 2015.

- [38] Saiful Hoque, Mathias Santos De Brito, Alexander Willner, Oliver Keil, and Thomas Magedanz. Towards container orchestration in fog computing infrastructures. In *2017 IEEE 41st Annual Computer Software and Applications Conference (COMPSAC)*, volume 2, pages 294–299. IEEE, 2017.
- [39] Falk Howar, Malte Isberner, Maik Merten, and Bernhard Steffen. Learnlib tutorial: from finite automata to register interface programs. In *International Symposium On Leveraging Applications of Formal Methods, Verification and Validation*, pages 587–590. Springer, 2012.
- [40] Falk Howar and Bernhard Steffen. Active automata learning in practice. In *Machine Learning for Dynamic Software Analysis: Potentials and Limits*, pages 123–148. Springer, 2018.
- [41] Malte Isberner, Falk Howar, and Bernhard Steffen. The open-source learnlib. In *International Conference on Computer Aided Verification*, pages 487–495. Springer, 2015.
- [42] Ramon Janssen, Frits W Vaandrager, and Sicco Verwer. Learning a state diagram of tcp using abstraction. *Bachelor thesis, ICIS, Radboud University Nijmegen*, 12, 2013.
- [43] Vineet John and Xia Liu. A survey of distributed message broker queues. *arXiv preprint arXiv:1704.00411*, 2017.
- [44] Asif Khan. Key characteristics of a container orchestration platform to enable a modern application. *IEEE cloud Computing*, 4(5):42–48, 2017.
- [45] Paul Menage. Control groups definition, implementation details, examples and api, 2004.
- [46] Maik Merten. Active automata learning for real life applications. 2013.
- [47] Maik Merten, Bernhard Steffen, Falk Howar, and Tiziana Margaria. Next generation learnlib. In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, pages 220–223. Springer, 2011.
- [48] Thanh-Tung Nguyen, Yu-Jin Yeom, Taehong Kim, Dae-Heon Park, and Sehan Kim. Horizontal pod autoscaling in kubernetes for elastic container orchestration. *Sensors*, 20(16):4621, 2020.
- [49] Claus Pahl and Brian Lee. Containers and clusters for edge cloud architectures—a technology review. In *2015 3rd international conference on future internet of things and cloud*, pages 379–386. IEEE, 2015.

- [50] Harald Raffelt, Bernhard Steffen, and Therese Berg. Learnlib: A library for automata learning and experimentation. In *Proceedings of the 10th international workshop on Formal methods for industrial critical systems*, pages 62–71, 2005.
- [51] Nathan Regola and Jean-Christophe Ducom. Recommendations for virtualization technologies in high performance computing. In *2010 IEEE Second International Conference on Cloud Computing Technology and Science*, pages 409–416. IEEE, 2010.
- [52] Fernando Rodríguez-Haro, Felix Freitag, Leandro Navarro, Efraín Hernández-sánchez, Nicandro Farías-Mendoza, Juan Antonio Guerrero-Ibáñez, and Apolinar González-Potes. A summary of virtualization techniques. *Procedia Technology*, 3:267–272, 2012.
- [53] Robert Rose. Survey of system virtualization techniques. 2004.
- [54] Maciej Rostanski, Krzysztof Grochla, and Aleksander Seman. Evaluation of highly available and fault-tolerant middleware clustered architectures using rabbitmq. In *2014 federated conference on computer science and information systems*, pages 879–884. IEEE, 2014.
- [55] Sasidhar Sekar. Autoscaling in kubernetes: Why doesn’t the horizontal pod autoscaler work for me? <https://medium.com/expedia-group-tech/autoscaling-in-kubernetes-why-doesnt-the-horizontal-pod-autoscaler-work-for-me-5f00> 2020. Accessed: 2022-09-29.
- [56] Wouter Smeenk. Applying automata learning to complex industrial software. *Master’s thesis, Radboud University Nijmegen*, 2012.
- [57] Rick Smetsers, Michele Volpato, Frits Vaandrager, and Sicco Verwer. Bigger is not always better: on the quality of hypotheses in active automata learning. In *International Conference on Grammatical Inference*, pages 167–181. PMLR, 2014.
- [58] Stephen Soltesz, Herbert Pötl, Marc E Fluczynski, Andy Bavier, and Larry Peterson. Container-based operating system virtualization: a scalable, high-performance alternative to hypervisors. In *Proceedings of the 2Nd ACM SIGOPS/EuroSys european conference on computer systems 2007*, pages 275–287, 2007.
- [59] Bernhard Steffen, Falk Howar, and Malte Isberner. Active automata learning: from dfas to interface programs and beyond. In *International Conference on Grammatical Inference*, pages 195–209. PMLR, 2012.
- [60] Bernhard Steffen, Falk Howar, and Maik Merten. Introduction to active automata learning from a practical perspective. In *International School on Formal Methods*

- for the Design of Computer, Communication and Software Systems*, pages 256–296. Springer, 2011.
- [61] Malte Isberner¹ Falk Howar² Bernhard Steffen. Learnlib: An open-source framework for active automata learning.
 - [62] Martin Tappler, Bernhard K Aichernig, and Roderick Bloem. Model-based testing of communication via active automata learning. In *2017 IEEE International conference on software testing, verification and validation (ICST)*, pages 276–287. IEEE, 2017.
 - [63] Max Tijssen, Erik Poll, and Joeri de Ruiter. Automatic modeling of ssh implementations with state machine learning algorithms. *Bachelor thesis, Radboud University, Nijmegen*, 2014.
 - [64] Eddy Truyen, Dimitri Van Landuyt, Davy Preuveneers, Bert Lagaisse, and Wouter Joosen. A comprehensive feature comparison study of open-source container orchestration frameworks. *Applied Sciences*, 9(5):931, 2019.
 - [65] James Turnbull. *The Docker Book: Containerization is the new virtualization*. James Turnbull, 2014.
 - [66] Mark Utting and Bruno Legeard. *Practical model-based testing: a tools approach*. Elsevier, 2010.
 - [67] Bill Venners. The java virtual machine. *Java and the Java virtual machine: definition, verification, validation*, 1998.
 - [68] W Eric Wong, Vidroha Debroy, Adithya Surampudi, HyeonJeong Kim, and Michael F Siok. Recent catastrophic accidents: Investigating how software was responsible. In *2010 Fourth International Conference on Secure Software Integration and Reliability Improvement*, pages 14–22. IEEE, 2010.
 - [69] Miguel G Xavier, Marcelo V Neves, Fabio D Rossi, Tiago C Ferreto, Timoteo Lange, and Cesar AF De Rose. Performance evaluation of container-based virtualization for high performance computing environments. In *2013 21st Euromicro International Conference on Parallel, Distributed, and Network-Based Processing*, pages 233–240. IEEE, 2013.

Hiermit versichere ich, dass ich die vorliegende Arbeit selbstständig verfasst habe und keine anderen als die angegebenen Quellen und Hilfsmittel verwendet sowie Zitate kenntlich gemacht habe.

Dortmund, den November 1, 2022

Julien Saevecke

